



Logistic Regression - Day 1



Agenda

Logistic Regression - Day 1

In this session, you will be learning:

01

**Classification
problems**

02

**Solving
classification
problems using
linear
regression**

03

**Interpreting
coefficients of
a logistic model**

04

**Finding out the
optimal
coefficients**

05

Python demo

Classification Problems

A grayscale photograph of five people (three men and two women) sitting together and smiling. The image is partially obscured by a dark gray overlay on the left and a solid red bar at the bottom. A white line graphic, consisting of a series of connected rounded and straight segments, is drawn over the right side of the image, starting from the middle and extending towards the bottom right corner.



Classification Problems

Predict if a machine will fail in the next 14 days.

VibrationX_14day	VibrationY_14day	VibrationZ_14day	Failed
...	Yes
...	No





Classification Problems

The HR department of a company wants to understand which employees are at risk of resigning.

# promotions	Current salary	Market Salary	Resigned
...	Yes
...	No





Classification Problems

Can we predict which patients are at risk of re-admission?

Patient ID	Age	Gender
001	23	M	...
Patient ID	Age	Gender





Classification Problems: Class Exercise

Take five minutes and discuss two scenarios in which the prediction problem is a classification problem. Also, discuss what kind of data you will need to collect.

- **Customer Churn Prediction:** Demographics, Past association with a product, and Number of times complaints registered
- **Predicting Fraud:** Demographics, Financial History, and Circumstances of transaction




Fitting Lines





Supervised Learning: Fitting lines




Feature 1 (Age_Normalized)	Feature 2 (Income_Normalized)	Response (Good/Bad)	Feature 1 (Age_Normalized)
.....	Good = 1
.....	Bad = 0

Fit a model of the form, $\text{Response} = b_0 + b_1 \text{Feature}_1$



Supervised Learning: Fitting lines



We fit a straight line to the data, where the response is binary in nature $y \in \{0, 1\}$.

Notice the predictions in red color

1. Predictions < 0 or Predictions > 1
2. $0 < \text{Predictions} < 1$

We are trying to fit:

$$\hat{y} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon = f(X_1, X_2)$$

The problem with the fitting linear model is:

$$f(X_1, X_2) \in (-\infty, +\infty)$$

$$y \in \{0, 1\}$$



Supervised Learning: Fitting lines


Age	Target
20	1
20	0
21	1
21	0



Age	Proportion (1)
20	0.50
21	0.50



Supervised Learning: Fitting lines



Instead of estimating $y \in \{0,1\}$ we can try to estimate the $\text{Prob}(y=1) = \hat{p}$

$$\hat{p} \in (0,1)$$

These estimates make sense now.

We are trying to fit:

$$\hat{p} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon = f(X_1, X_2)$$

The problem still is:

$$f(X_1, X_2) \in (-\infty, +\infty)$$

$$\hat{p} \in (0,1)$$



Supervised Learning: Fitting lines

Age	Target
20	1
20	1
20	1
20	0
21	1
21	0



Age	Pr (1)	Pr/1-Pr
20	0.75	3
21	0.50	0.5



Supervised Learning: Fitting lines

Age	Target
20	1
20	1
20	1
20	0
21	1
21	0

Age	Pr (1)	Pr/1-Pr
20	0.75	3
21	0.50	0.5

Maximum and Minimum value $\text{Pr}/(1-\text{Pr})$?

Minimum value of $\text{Pr}(1)$? $\text{Pr}(1) = 0$

Maximum value of $\text{Pr}(1)$? $\text{Pr}(1) = 1$


Min $\text{Pr}(1) = 0$, $\text{Pr}/(1-\text{Pr}) = ?$ $\text{Pr}/(1-\text{Pr}) = 0$

Max $\text{Pr}(1) = 1$, $\text{Pr}/(1-\text{Pr}) = ?$ $\text{Pr}/(1-\text{Pr}) = \text{Infinity}$

$$0 \leq \text{Pr}/(1-\text{Pr}) \leq \text{Infinity}$$



Supervised Learning: Fitting lines



Instead of estimating $y \in \{0,1\}$ we can try to estimate the $\hat{p} / (1 - \hat{p})$

$$\hat{p} / (1 - \hat{p}) \in (0, +\infty)$$

These estimates make sense now.

We are trying to fit:

$$\hat{p} / (1 - \hat{p}) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \varepsilon = f(X_1, X_2)$$

The problem still is:

$$f(X_1, X_2) \in (-\infty, +\infty)$$

$$\hat{p} / (1 - \hat{p}) \in (0, +\infty)$$



Supervised Learning: Fitting lines

Age	Target
20	1
20	1
20	1
20	0
21	1
21	0




Age	Pr (1)	Pr/1-Pr	log(Pr/(1-Pr))
20	0.75	3	1.09
21	0.50	0.5	-0.70

$$0 \leq \text{Pr}/(1-\text{Pr}) \leq \text{Infinity}$$

$$-\text{Infinity} \leq \log(\text{Pr}/(1-\text{Pr})) \leq \text{Infinity}$$



Supervised Learning: Fitting lines



Instead of estimating $y \in \{0,1\}$ we can try to estimate the $\log(\hat{p} / (1 - \hat{p}))$

$$\log \hat{p} / (1 - \hat{p}) \in (-\infty, +\infty)$$

All these estimates make sense now.

We are trying to fit:

$$\log \hat{p} / (1 - \hat{p}) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \varepsilon = f(X_1, X_2)$$

$$f(X_1, X_2) \in (-\infty, +\infty) \log \hat{p} / (1 - \hat{p}) \in (-\infty, +\infty)$$




Supervised Learning: Fitting lines

$$\log \frac{\hat{p}}{1-\hat{p}} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \varepsilon$$

$$\hat{p} = \frac{e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \varepsilon}}{1 + e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \varepsilon}}$$

Logistic Function





Classification Problems: Class Exercise

Imagine that there are 125 customers of age 25 years. Of them, 25 have subscribed to a premium subscription of an OTT platform. Find out the odds ratio.

$$p = \frac{25}{125}, 1 - p = \frac{100}{125}, \text{so odds ratio} = \frac{p}{1 - p} = \frac{25}{125} * \frac{125}{100} = \frac{1}{4}$$





Interpreting Coefficients of a Logistic Model




Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log(p/(1-p)) = 2.1 + 0.08 \text{Age}$$

Age	Churned
28	Yes
32	Yes
40	No





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log(p/(1-p)) = 2.1 + 0.08 \text{Age}$$



Changing Age by 1 unit
will change the log odds
of someone churning by
0.08 units.





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log\left(\frac{p}{1-p}\right) = 2.1 + 0.08 \text{Age}$$



Changing Age by 1 unit
will change the log odds
of someone churning by
0.08 units.





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log(p/(1-p)) = 2.1 + 0.08 * 20$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log(p/(1-p)) = 2.1 + 1.6$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log(p/(1-p)) = 3.7$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log\left(\frac{p}{1-p}\right) = e^{3.7}$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log\left(\frac{p}{1-p}\right) = 2.71^{3.7}$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$\log\left(\frac{p}{1-p}\right) = 39.99$$





Interpreting Logistic Regression Coefficients

Predicting churn based on a person's age

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1$$

$$P = 0.97$$





Class Exercise

- Imagine you were trying to model the propensity of customers to churn, and you were able to build the following logistic regression model:

$$\log \frac{p}{1-p} = 1.95 - 2.5Age + 1.68Income$$

- Can we conclude that as age increases, the propensity for a customer to churn decreases, keeping all else constant?
- What would be the churn probability for a person aged 25 with an income of 15000?





Class Exercise

- Imagine you were trying to model the propensity of customers to churn, and you were able to build the following logistic regression model:

$$\log \frac{p}{1-p} = 0.0001 - 0.005Age + 0.008Income$$

- Can we conclude that as age increases, the propensity for a customer to churn decreases, keeping all else constant? (Yes, negative sign on the coefficient of Age.)
- What would be the probability of churn for a person with age 25 and income 150? ($p = \frac{E}{1+E}$, where $E = e^{0.0001-0.005*25+0.008*150}$, $p = 0.75$)





Finding Out the Optimal Coefficients



Estimating Coefficients of Logistic Model

Age	Good_Bad (Good =1)	Prediction
20	1	
21	1	
24	0	
25	0	
29	0	
30	1	
38	1	

$$\beta_0 = 0.7$$
$$\beta_1 = 1.7$$

$$\log \frac{\hat{p}}{1-\hat{p}} = \beta_0 + \beta_1 \text{Age}$$

$$\hat{p} = \frac{e^{\beta_0 + \beta_1 \text{Age}}}{1 + e^{\beta_0 + \beta_1 \text{Age}}}$$

$$\hat{p} = \frac{e^{0.7 + 1.7 \text{Age}}}{1 + e^{0.7 + 1.7 \text{Age}}}$$



Estimating Coefficients of Logistic Model

Age	Good_Bad (Good =1)	Prediction
20	1	
21	1	
24	0	
25	0	
29	0	
30	1	
38	1	

$$\beta_0 = 0.7$$
$$\beta_1 = 1.7$$

$$\log \frac{\hat{p}}{1-\hat{p}} = \beta_0 + \beta_1 \text{Age}$$

$$\hat{p} = \frac{e^{\beta_0 + \beta_1 \text{Age}}}{1 + e^{\beta_0 + \beta_1 \text{Age}}}$$

Age	Good_Bad (Good =1)	Prediction
20	1	
21	1	
24	0	
25	0	
29	0	
30	1	
38	1	

$$\beta_0 = 0.3$$
$$\beta_1 = 2.2$$



Estimating Coefficients of Logistic Model

Age	Good_Bad (Good =1)	Prediction
20	1	
21	1	
24	0	
25	0	
29	0	
30	1	
38	1	

$$\beta_0 = 0.7$$
$$\beta_1 = 1.7$$

$$\log \frac{\hat{p}}{1-\hat{p}} = \beta_0 + \beta_1 \text{Age}$$

$$\hat{p} = \frac{e^{\beta_0 + \beta_1 \text{Age}}}{1 + e^{\beta_0 + \beta_1 \text{Age}}}$$

Age	Good_Bad (Good =1)	Prediction
20	1	0.70
21	1	0.60
24	0	0.50
25	0	0.45
29	0	0.70
30	1	0.62
38	1	0.40

$$\beta_0 = 0.3$$
$$\beta_1 = 2.2$$



Estimating Coefficients of Logistic Model

Age	Good_Bad (Good =1)	Prediction
20	1	
21	1	
24	0	
25	0	
29	0	
30	1	
38	1	

$$\beta_0 = 0.7$$
$$\beta_1 = 1.7$$

Clearly $\beta_0 = 0.7$ and $\beta_1 = 1.7$ is a better choice

$$\log \frac{\hat{p}}{1-\hat{p}} = \beta_0 + \beta_1 \text{Age}$$

$$\hat{p} = \frac{e^{\beta_0 + \beta_1 \text{Age}}}{1 + e^{\beta_0 + \beta_1 \text{Age}}}$$

Age	Good_Bad (Good =1)	Prediction
20	1	0.70
21	1	0.60
24	0	0.50
25	0	0.45
29	0	0.70
30	1	0.62
38	1	0.40

$$\beta_0 = 0.3$$
$$\beta_1 = 2.2$$



Estimating Coefficients of Logistic Model

- One would like to choose the model coefficients so that the model gives a high score to events and a low score to non-events.
- But how will we measure a model's ability to assign a high score to events and a low to non-events? Cost Function.

See the excel sheet **logistic_cost.xlsx**





Thank You!

Copyright © HeroX Private Limited, 2023. All rights reserved.